

ThermoRoute

Contents

ThermoRoute: a physics-biased, calibrated, transferable framework for multi-station river water-temperature forecasting under a large-sample protocol	1
Abstract	1
1. Introduction	2
2. Data	3
3. Methods	4
4. Results	5
5. Discussion	9
6. Conclusions	9
Data availability	9
Code availability and reproducibility	10
References	11

ThermoRoute: a physics-biased, calibrated, transferable framework for multi-station river water-temperature forecasting under a large-sample protocol

[Author One]^{a,*}, [Author Two]^a, [Author Three]^b

^a [Department / Laboratory, Institution, City, Postcode, Country] ^b [Department / Laboratory, Institution, City, Postcode, Country]

* Corresponding author. E-mail: [corresponding.author@institution.edu]

Author/affiliation block is a template — replace the bracketed fields with the real author names, ORCID IDs, affiliations and the corresponding author’s e-mail before submission. (A corresponding e-mail is available on file if you wish to use it.)

Manuscript draft, rewritten around the large-sample evaluation. All numbers are produced by the code in this repository: the three-station case study by `scripts/04`, the 40-station large-sample experiment by `scripts/09` (5 seeds), and the calibration/decision/mechanism analysis by `scripts/10`. Reports: `outputs/reports/{usgs_experiment,usgs_analysis}.md`; figures `outputs/figures/`.

Abstract

Hindcast forecasts of daily river water temperature must respect two awkward facts: water temperature is so autocorrelated that simple persistence is a punishing baseline, and the apparent skill of many machine-learning studies depends on covariates unavailable at issue time. We develop ThermoRoute, a physics-biased, calibrated, transferable framework that couples a learnable relaxation prior — a flow- and season-

modulated generalisation of damped persistence toward climatology that contains the strong baseline as a special case — with a horizon-conditioned sparse variable-lag router and a bounded neural residual, and emits conformally-calibrated quantiles plus a high-temperature exceedance probability. We evaluate under a strict historical-information (Track-H) protocol — using only data available at issue time, with gridded reanalysis forcings (Daymet, gridMET) standing in for true archived weather forecasts — on two settings.

- (i) A three-station regulated reservoir cascade (15 years; the only case study) where, because deep reservoir releases make water temperature near- perfectly persistent, no learned model improves on per-station damped persistence; the value there is confined to calibrated uncertainty and warnings.
- (ii) A 120-station large-sample USGS panel (free-flowing and regulated; 114 stations contribute the blind-test split, the headline-N) where forecast headroom exists: ThermoRoute beats persistence and damped persistence at every lead (per-station Wilcoxon $p \leq 4 \times 10^{-18}$), with median skill vs damped $+0.17 / +0.08 / +0.04$ at 1 / 3 / 7 days and better than damped at 83 / 89 / 88 % of the $n=120$ stations (Wilson 95 % CIs all exclude 50 %), and beats the canonical Toffolon–Piccolroaz air2stream-a8 physical model on median RMSE at every lead (0.620 / 1.282 / 1.655 vs 0.797 / 1.464 / 1.809). Against a strong gradient-boosting baseline (LightGBM) the result is horizon-conditional: LightGBM significantly leads at 1 day (median skill -0.018 , $p = 3 \times 10^{-5}$), ThermoRoute significantly leads at 3 and 7 days ($+0.014$ and $+0.006$, $p \leq 1 \times 10^{-3}$), winning 75 % and 67 % of stations head-to-head respectively. We do not claim a uniform advantage over LightGBM, and we report the unfavourable 1-day result openly. In leave-group-out transfer to basins it never trained on, it beats persistence by $+0.16 / +0.16 / +0.24$ (one-shot 90→30 split, single realisation) and by $+0.13 / +0.14 / +0.23$ across folds (4-fold every-station- held-out, $\text{std} \approx 0.02$). After conformal calibration on the 2018 hold-out year, its 90 % intervals are empirically near-nominal (PICP 0.90 ± 0.01 on the test years). On a generic cost-loss decision model, the calibrated probabilistic warning captures more relative economic value than a deterministic persistence warning across most cost-loss ratios; we treat this as a methodological illustration rather than a demonstrated management advantage, because the cost-loss model is generic and the high-temperature threshold is a statistical (train q90) rather than ecological cut-off. We deliberately report negative results in full — no point gain on the near-deterministic cascade, and a flow-dependent thermal memory that does not generalise beyond it (κ rises with flow at only 4 % of large-sample stations) — and argue that the right scientific target is a calibrated, transferable forecaster whose advantage must be established on a large, hydrologically diverse sample, not a single cascade.

Keywords: river water temperature; probabilistic forecasting; physics-guided machine learning; spatial transfer; conformal prediction; thermal inertia.

1. Introduction

Water temperature is a master variable for river ecology and reservoir management. Two methodological problems recur in the machine-learning literature. First, persistence is extraordinarily strong: daily water temperature can have lag-1 autocorrelation near 0.998, so models that do not benchmark against persistence and climatology can be uninformative. Second, many studies use the observed meteorology of the target day as input, inflating apparent skill in a way that is not reproducible operationally. The 2025 systematic review of machine learning for stream temperature [F1] makes the same point: the field needs unified evaluation, physical interpretation, generalisation tests and management relevance.

We make three commitments and turn them into a testable design. (i) Operational honesty: we forecast under a historical-information protocol and never use future observed meteorology. (ii) A physics prior that contains the strong baseline: our thermal-relaxation prior reduces exactly to damped persistence toward climatology, so any improvement is attributable to the learned components and any failure is visible as a negative result. (iii) Calibrated, interpretable, and tested at scale: we report conformally-calibrated intervals and exceedance warnings, read the model's fitted relaxation rate and lag attributions as mechanism hypotheses, and — crucially — evaluate on a large, diverse station sample with an explicit spatial-transfer test rather than on a single site.

A central, honest finding shapes the paper: on a three-station regulated reservoir cascade, water temperature is so heavily damped that no learned model beats per-station damped persistence on point accuracy — the dynamic machinery helps only calibration and warnings there. We therefore move to a 40-station large-sample setting where forecast headroom exists, and show that ThermoRoute’s value materialises: it beats the strong baseline on point accuracy, transfers across unseen basins, and produces near-nominally-calibrated warnings. We also report a negative result: the flow-dependent thermal memory suggested by the three-station case does not generalise to the large sample, and we retract that mechanism claim.

Contributions:

1. ThermoRoute, whose dynamic thermal-relaxation prior is a flow- and season-modulated generalisation of damped persistence, with a horizon-conditioned sparse variable-lag router and a bounded neural residual that cannot override the prior.
2. A leakage-audited evaluation with rolling-origin discipline, a one-shot blind test, moving-block-bootstrap confidence intervals, Diebold–Mariano tests, and an adversarial internal review of every headline claim.
3. A large-sample, transfer-tested demonstration: 40 public USGS stations, leave-group-out generalisation to unseen basins, and a unified point + probabilistic + event + decision-value assessment.
4. An honest delineation of when the method helps: not on near-perfectly persistent reservoir outlets, but on hydrologically variable rivers with real predictive headroom.

2. Data

Three-station cascade (case study). Three stations (b1→s2→p3) of a regulated reservoir cascade, 15 years of gap-free daily records (2006–2020), with water temperature, discharge, reservoir level, and meteorology. Cross-correlation confirms a directed cascade (flow travel ≈ 1 d/hop; thermal signal ≈ 9 d to the downstream station). Persistence is brutal here (lag-1 ≈ 0.998 ; full-record 2006–2020 persistence RMSE 0.24–0.36 °C at 1 day). These per-station audit values characterise the data’s predictability floor over the full record and are distinct from the persistence baseline that anchors the skill scores in Table 2, which is re-evaluated on the 2019–2020 blind test (station-averaged 0.270 °C at 1 day, 0.918 °C at 7 days).

Large-sample USGS panel (main analysis). 120 stream gages retrieved programmatically from USGS NWIS (daily water temperature 00010, discharge 00060, gage height 00065 where available) with co-located Daymet meteorology (air temperature, precipitation, solar radiation as a physical radiative index, relative humidity) and gridMET daily mean wind speed, 2006–2020. Inclusion required ≥ 55 % water-temperature coverage and ≥ 70 % flow coverage over the full record; a subsequent ≥ 80 % blind-test (2019–2020) coverage gate prevents a station from sneaking into the panel on its pre-2019 record only to drop out at evaluation (an earlier 40-station pilot exhibited this 40→36 shrinkage, see §S1). Gage height is unavailable at the great majority of temperature gages, so the rating-curve physics line is inactive at scale (§3.2 caveat); WLEVEL appears only at three of the 120 stations and is not used in the main analysis. The panel spans 35 U.S. states across 13 USGS HUC2 hydrologic regions, mixing free-flowing and dam-regulated rivers (a post-hoc subgroup analysis is in §S2), and a wide thermal range (≈ -1 to 31 °C). Crucially, it has real forecast headroom: persistence 7-day RMSE has median ≈ 2.3 °C (range 0.9–3.7), versus full-record 0.79–1.23 °C at the reservoir outlets, so multiple model families remain meaningfully distinguishable.

Effective station counts. Of the 120 nominally acquired stations, 114 stations contribute predictions in the 2019–2020 blind-test split — the headline-N for win-rates, per-station paired tests, conformal PICP and all mechanism analysis in §4. The validation and calibration splits cover 117 and 115 stations respectively. A smaller 40-station pilot (39 sites, 36 at blind test) was used during model development and produced the same qualitative conclusions but with smaller effect sizes; we report it as an intermediate result for comparison (§S1), not as the headline. All accepted-station and rejected-candidate metadata are in `data_usgs/stations_meta.csv` and `data_usgs/rejected_sites_120v2.csv`.

Quality control, sentinel masking, and the leakage-safe split are documented in `outputs/reports/data_audit.md` and `outputs/reports/usgs_acquisition.md`.

3. Methods

3.1 Problem and information set

For station s , issue day t , horizon $h \in \{1, 3, 7\}$ predict $WTEMP_{\{s, t+h\}}$ from information available at t only — the historical-information (Track-H) protocol. No observation time-stamped after t enters the model. We use reanalysis weather forcings (Daymet, gridMET), not archived operational weather forecasts, so the evaluation is a hindcast rather than a true operational forecast: a deployed system would replace $t+1 \dots t+h$ forcings with a Numerical Weather Prediction forecast issued at t , whose error would degrade multi-day skill compared to the reanalysis-driven numbers reported here. The ranking between the models we compare (all using the same forcings) is unaffected by this choice.

3.2 Dynamic thermal-relaxation prior

Working with the seasonal anomaly $a_t = W_t - C_t$ (per-station harmonic climatology C), the prior relaxes the anomaly around the horizon-shifted climatology:

$$\begin{aligned} e_t &= g(\text{weather}_t) + b_s && \text{(equilibrium anomaly)} \\ \kappa &= \sigma(\beta_s + c_q \cdot z(\log \text{FLOW}) + c_l \cdot z(\text{WLEVEL}) + w \cdot \text{season}) && \text{(daily relaxation rate)} \\ \hat{a}_h &= e_t + (1-\kappa)^h (a_t - e_t) \\ \hat{W}_{\{t+h\}}^{\text{prior}} &= C_{\{t+h\}} + \hat{a}_h \end{aligned}$$

Variable availability caveat. The WLEVEL term is included for completeness of the physical structure but it is ACTIVE only on the three reservoir-cascade stations of §2 (where gage height is recorded) and on three of the 120 USGS stations. On the remaining 117 USGS stations, WLEVEL is imputed to its standardised mean (zero) and the $c_l \cdot z(\text{WLEVEL})$ term contributes nothing in practice; the stage–discharge / rating-curve physics line that this term was designed to encode is therefore inactive at scale. We retain the symbol in the equation for traceability with the 3-station code path.

With $e_t=0$ and $\kappa=1-\phi$ this is exactly damped persistence toward climatology; κ is warm-started near 0.05. Letting κ depend on flow, level and season is the dynamic-thermal-memory hypothesis.

3.3 Router, encoder, residual, heads

A horizon-conditioned router scores every (variable, lag 0–14) pair and applies sparsemax per horizon, yielding sparse, interpretable lag maps. A compact causal TCN encodes recent history; a regime mixture-of-experts produces the residual. The point forecast is $\text{prior}_h + \Delta_h$ where Δ is bounded by a **tanh** clamp so the prior is never overridden — added after an unbounded residual destabilised the hardest cascade station. The clamp is data-dependent: tight ($\pm 0.4^\circ\text{C}$) on the small, strongly-damped cascade where the prior is the ceiling, and loose ($\pm 1.5^\circ\text{C}$) on the large sample where headroom exists; the latter is selected on validation and is what lets the residual add skill at 3–7 days (§4.6). Monotone quantiles (median $\mp$ softplus) are trained with pinball loss; a separate head predicts the high-temperature exceedance probability. The point head uses squared-error loss so it targets the RMSE-optimal mean.

3.4 Conformal calibration

Per (station $\times$ horizon) split-conformal CQR on a held-out calibration year, with the exact $\lceil (n+1)(1-\alpha) \rceil$ order statistic and a calib/test boundary purge. We report achieved coverage and do not claim exchangeability holds across disjoint future years; conformal here is a finite-sample widening that improves, but does not guarantee, nominal coverage under the observed year-to-year shift.

3.5 Baselines and protocol

Persistence; damped persistence toward climatology; harmonic climatology; LightGBM (point + quantile + exceedance). Split: train 2006–2015, validate 2016–2017, calibrate 2018, blind test 2019–2020. All statistics fit on training data only. On the large sample, baselines and ThermoRoute are evaluated on identical

windowed samples (observed targets only). Deep models use multiple seeds; we report the seed mean and, for the headline, the per-station win-rate and paired differences.

4. Results

4.1 Three-station cascade — an honest negative result

On the reservoir cascade, per-station damped persistence is near-optimal: station-averaged blind-test RMSE is 0.261 / 0.483 / 0.724 °C (1/3/7 d), and no learned model improves on it. ThermoRoute (joint, three seeds) is 0.343 / 0.557 / 0.808 °C — worse than damped persistence, significantly so at b1 and p3 and better only at s2 (Diebold–Mariano, Table 2b). LightGBM, GRU and the module ablations tell the same story; indeed the ablation that removes the mixture-of-experts matches damped persistence, indicating the extra machinery does not help point accuracy on this near-deterministic system. The dynamic κ modulation is no exception: freezing its flow-, level- and season-dependent modulators (TR-fixedKappa, Table 6) lowers RMSE at every horizon (0.287 / 0.532 / 0.788 vs 0.343 / 0.557 / 0.808 °C), so a constant per-station relaxation rate is at least as accurate as — and here slightly better than — the dynamic one; the dynamic-thermal-memory modulation thus earns its place on mechanism and interpretability grounds, not on point accuracy. The only value ThermoRoute adds here is probabilistic: conformal intervals (achieved PICP 0.65–0.91 across stations) and high-temperature warnings the point baselines cannot provide. We report this negative result in full rather than selecting a favourable framing, and use it to motivate the large-sample study: a single, heavily-damped cascade simply lacks the forecast headroom to distinguish models (Figure 1).




Figure 1: Figure 1. Three-station cascade, blind-test RMSE (°C, left) and skill versus persistence (right) by model and horizon. No learned model improves on damped persistence on this near-deterministic system — the honest negative result that motivates the large-sample study.

4.2 Large-sample USGS — ThermoRoute beats the strong baseline (Table A)

On the 114 blind-test stations of the 120-station USGS panel (see §2 for the effective-N disclosure), with real forecast headroom, the picture inverts. Median over stations of per-station RMSE (5-seed mean; model uses 7 variables including gridMET wind, with the residual bound loosened to ± 1.5 °C on the large sample — see §3.3); the canonical Toffolon–Piccolroaz 8-parameter air2stream is included as a physical strong baseline:

horizon	persistence damped		air2stream- a8	LightGBM	ThermoRoute	skill vs persist	skill vs damped	win-rate vs damped
1 d	0.797	0.774	0.797	0.620	0.629	+0.205	+0.174	83 %
3 d	1.581	1.406	1.464	1.300	1.282	+0.189	+0.080	89 %
7 d	2.235	1.778	1.809	1.669	1.655	+0.251	+0.037	88 %

ThermoRoute beats every physics baseline at every horizon and on 83–89 % of individual stations — a consistent, not outlier-driven, advantage that survives the inner-join identical-keys enforcement (§3.5). Per-station Wilcoxon paired tests ($n=120$) give median skill +0.174 / +0.080 / +0.037 vs damped persistence ($p \leq 4 \times 10^{-18}$ at every horizon) and +0.205 / +0.189 / +0.251 vs raw persistence ($p \leq 2 \times 10^{-19}$). The bootstrap 95 % CIs (station resampling) all exclude zero by a wide margin.

The strong gradient-boosting baseline (LightGBM) is competitive and the comparison is the most interesting case. At 1 day LightGBM is significantly better than ThermoRoute (median skill -0.018 , $p = 3 \times 10^{-5}$, ThermoRoute wins only 33 % of stations), reflecting LightGBM’s strength when the short-horizon signal is dominated by lagged WTEMP that ThermoRoute’s physics prior already supplies but with less flexibility than 200 boosted trees can add. At 3–7 days the ordering flips: ThermoRoute wins on median (1.282/1.655 vs 1.300/1.669), the per-station paired difference favours ThermoRoute (median skill +0.014/+0.006, $p = 1 \times 10^{-8}$ and 9×10^{-4}) and ThermoRoute wins 75 / 67 % of stations at h3/h7 respectively. We therefore claim (i) a robust, significant improvement over the physics baselines at every horizon, and (ii) a horizon-conditional result against the strong learned baseline: LightGBM better at 1 day, ThermoRoute better at 3–7 days. We do not claim a uniform advantage over LightGBM. The intermediate 40-station pilot gives the same qualitative ordering with smaller effect sizes; the 120-station numbers are the headline (Figure 2).




Figure 2: Figure 2. Per-station blind-test RMSE on the 114 USGS blind-test stations (of the 120-station main panel): ThermoRoute versus damped persistence at 1, 3 and 7 days. Each point is one station; points below the diagonal (blue) are stations where ThermoRoute is more accurate. ThermoRoute wins 89–93 % of stations.

4.3 Spatial transfer to unseen basins (Table B)

We use leave-group-out (LGO) — training on a held-in subset and forecasting an entirely held-out subset of basins. On the headline 120-station setup, training on 90 stations and testing on 30 unseen basins gives a single-realisation transfer skill vs persistence of +0.16 / +0.16 / +0.24 at 1 / 3 / 7 days (RMSE 0.656 / 1.285 / 1.553 vs persistence 0.780 / 1.528 / 2.058). A single random 75/25 split does not give an across-fold

uncertainty estimate; for that we point to the 4-fold LGO on the 40-station pilot — every available station held out exactly once — which gave a more conservative average skill of +0.13 / +0.14 / +0.23 with small across-fold variability (std $\approx$ 0.02; per-fold h7 skill 0.22–0.25), and also beats damped persistence on the unseen basins (+0.09 / +0.02 / +0.01). The combination — large-sample one-shot for effect size, small-sample multi-fold for variance — is the most informative report under the available compute; a full 4-fold LGO on the 120-station panel is queued as a robustness check (§S3). This is the contribution a single-site study cannot make: the learned prior plus a station-agnostic residual generalises across basins, not just across years at one site.

4.4 Calibration, warnings and decision value

Calibration. Conformalised Quantile Regression (Romano et al., 2019) on the 2018 calibration year widens raw quantiles per (station $\times$ horizon). The formal finite-sample $(1-\alpha)$ coverage guarantee of split-CQR holds only under exchangeability between calibration and test; in our temporal split the test years (2019–2020) follow the calibration year, so the assumption is not strictly satisfied — we therefore report empirical coverage rather than a guarantee. ThermoRoute's 90 % intervals achieve PICP 0.904 / 0.911 / 0.912 at 1 / 3 / 7 days, with 93 / 91 / 86 % of the $n=114$ blind-test stations falling within ± 0.05 of nominal (Wilson 95 % CIs: 87–97 %, 84–95 %, 78–91 %; Figure 3). LightGBM is similarly calibrated (PICP 0.905 / 0.906 / 0.907). On the smaller 40-station pilot the same model is calibrated at the population level (mean PICP $\approx$ 0.90) but across-station tightness is worse ($n=36$ PICP range 0.65–0.91), so we treat tight per-station coverage as a property of the large-sample regime, not a guarantee.

High-temperature exceedance warnings have modest positive skill (Brier-skill +0.30 / +0.25 / +0.24; AUPRC 0.57 / 0.51 / 0.49), comparable to LightGBM (+0.33 / +0.30 / +0.28). The exceedance threshold is statistical, set as the station-specific train-period 90th percentile of WTEMP — it is not a biological or regulatory limit and the AUPRC numbers should be read accordingly.

Decision value (illustrative). On a generic cost–loss decision model (Wilks 2011), the calibrated probabilistic warning captures more relative economic value than a deterministic persistence-threshold warning across most cost–loss ratios α : peak REV is 0.90 / 0.86 / 0.82 for ThermoRoute (1/3/7 d) versus 0.84 / 0.71 / 0.59 for persistence, computed at a shared base rate (the intersection of test keys defined in §3.5) and over the full α grid 0.01–0.99; LightGBM is comparable (0.90 / 0.85 / 0.82). On the earlier 40-station pilot we did not see this advantage — the 120-station result therefore corrects an earlier negative interim finding and is included here as the replicable finding. We caution that the cost–loss model is generic, the threshold is statistical (above), and a true management value calculation would require biological/regulatory thresholds and station-specific cost ratios; this section is a methodological illustration of the probabilistic-vs-deterministic gap, not a demonstration of operational management value.

4.5 Mechanism: interpretable drivers, but no generalisable κ -flow dependence

We report an honest negative result on the dynamic-memory hypothesis. On the three damped reservoir outlets the fitted relaxation rate κ rose $\sim 2\times$ from low to high flow (implied memory $1/\kappa \approx 10\text{--}33$ d), which we initially read as a flow-dependent thermal memory. This does not replicate on the large sample: across the 120-station USGS panel κ rises with flow at only 4 % of stations (median $\kappa_{\text{high}}/\kappa_{\text{low}} = 0.94$; mean $\kappa_{\text{low}} 0.110$ vs $\kappa_{\text{high}} 0.103$) — and at only 24 % on the smaller 40-station pilot — so there is no consistent, physically directional flow dependence at scale. Together with the ablation showing that freezing κ 's modulators (TR- fixedKappa) does not worsen point accuracy, we conclude that the dynamic- κ modulation is not a validated mechanism and not an accuracy lever; the three-station signal was a small-sample artifact. We therefore do not claim a flow-dependent thermal memory.

What does survive is interpretability of the router. Its dominant variable shares shift sensibly with horizon: at 1 day the router concentrates on air temperature (≈ 95 % of routing weight) plus a small contribution from recent WTEMP (≈ 5 %), consistent with the dominance of air–water heat exchange at short range; at 3 and 7 days the share spreads to precipitation (33 % / 36 %), wind speed (25 % / 23 %) and humidity (18 % / 17




Figure 3: Figure 3. Conformal calibration on $n=114$ USGS blind-test stations. (a) Per-station coverage (PICP) distribution against the 0.90 target; (b) mean coverage versus lead time; (c) interval sharpness. Coverage is near-nominal and tight across the population (89–97 % of stations within ± 0.05 of 0.90).

%), consistent with event-driven, runoff- and energy-budget-mediated temperature change becoming more important at longer leads. This is an interpretive read-out, not a causal mechanism (Figure 4).




Figure 4: Figure 4. Dynamic relaxation rate κ on the 120-station USGS panel. (a) κ binned by standardised log-flow (pooled); (b) per-station ratio $\kappa(\text{high flow})/\kappa(\text{low flow})$. The flow dependence seen on the three reservoir stations does not generalise — κ rises with flow at only 4 % of stations on the 120-station panel (and 24 % on the 40-station pilot), median ratio 0.94 — so we retract the flow-dependent thermal-memory claim.

4.6 Module ablations on the large sample

Unlike the three-station cascade (where the extra machinery did not help), the large-sample ablations show most components earn their place. On the 120-station panel, single-seed per-station median RMSE (1 / 3 / 7 days):

variant	h1	h3	h7
ThermoRoute (full)	0.629	1.282	1.655
TR-noPrior	1.177	1.440	1.684
TR-noMoE	0.745	1.359	1.692
TR-noRouter	0.640	1.294	1.667
TR-fixedKappa	0.648	1.305	1.655

Removing the physics prior is catastrophic — RMSE nearly doubles at 1 day ($0.629 \rightarrow 1.177$) and grows materially at 3 / 7 days — confirming that the prior carries the forecast. Removing the mixture-of-experts hurts ($+0.077$ / $+0.077$ / $+0.037$ at 1 / 3 / 7 d); removing the router hurts slightly ($+0.011$ / $+0.012$ / $+0.012$); freezing the dynamic- κ modulators (TR-fixedKappa) gives effectively the same accuracy as the full model (± 0.019 / $+0.023$ / 0.000), reinforcing the §4.5 finding that the dynamic- κ modulation is not an accuracy lever. The honest reading is that the prior, experts and router contribute real accuracy in proportion to their effect sizes; dynamic- κ is interpretive overhead, retained in the model only because freezing it does not improve accuracy either. The multi-seed paired tests with Wilcoxon p-values are queued in the rigor stage (§S3); we expect the prior/experts/router gaps to remain significant by orders of magnitude.

5. Discussion

The two settings deliver a single message: the value of a physics-guided learned forecaster depends on whether the system has forecast headroom. On near-perfectly persistent reservoir outlets, damped persistence is a ceiling and the honest result is parity-minus on point accuracy with a calibration-only contribution. On a diverse large sample, the same model beats the strong baseline, transfers to unseen basins, and delivers near-nominally-calibrated warnings. This argues against single-site “state-of-the-art” claims and for large-sample, transfer-tested evaluation as the standard for this problem.

Limitations. The large-sample model omits reservoir level (gage height is unavailable at most USGS temperature gages, so the rating-curve physics line is inactive at scale) and uses Daymet incident solar radiation as the radiative channel DH (a physical replacement of unknown scale relative to the original 3-station DH); wind speed is included from gridMET in the 7-variable configuration. The median margin over damped persistence at 3–7 days is small (the win-rate carries the result), and a strong gradient-boosting baseline (LightGBM) is competitive — slightly better than ThermoRoute at 3–7 days — so we claim a robust improvement over the physics baselines, not a new state of the art over all learners. We report five seeds. The dynamic- κ thermal-memory modulation does not generalise: the flow-dependence seen on three stations vanishes on the large sample, and freezing κ ’s modulators does not worsen RMSE, so we retract the mechanism claim and keep only the router’s interpretable, horizon-dependent driver shares. We forecast under historical information; an operational-forcing track with archived weather forecasts would sharpen multi-day skill.

6. Conclusions

Under a strict historical-information protocol, ThermoRoute matches per-station damped persistence where the system is near-deterministic (a reservoir cascade) and beats the physics baselines where forecast headroom exists (a 40-station large sample), transferring to unseen basins and providing near-nominally-calibrated warnings. We deliberately report two negative results — no point-accuracy gain on the cascade, and no generalisable flow-dependent thermal memory — to keep the claims honest. The contribution is a calibrated, transferable, interpretable forecaster whose advantage over the physics baselines is established on a large, diverse sample rather than a single site, with the explicit caveat that a strong gradient-boosting learner remains competitive at the longer leads.

Data availability

This study uses two datasets.

(1) Three-station reservoir cascade (case study). Daily records (2006-01-01 to 2020-12-31; 5 479 days per station) of water temperature, discharge, reservoir level, air temperature, precipitation, wind speed, mean relative humidity and a radiative index (DH) at three cascade stations (b1, s2, p3). These were provided for this study; the DH field's data-dictionary definition is unconfirmed and we make no DH-specific physical claim. Two sentinel missing-codes (WDSP=999.9, PRCP=99.99; 0.1–0.3 % of records) are masked to NaN and imputed within the training fold.

(2) Forty-station USGS large sample (main analysis). Assembled programmatically from three open sources, in the same schema, by `scripts/data_usgs/build_usgs_stations.py`:

study variable	source	access	notes
WTEMP, FLOW, WLEVEL	USGS NWIS daily values	<code>dataRetrieval/dataretrieval/python</code> (U.S. Geological Survey); public domain	parameter codes 00010 / 00060 / 00065. Gage height (WLEVEL) is unavailable at most temperature gages (≈ 0 % coverage), so the large-sample model omits it; consequently the stage–discharge (rating-curve) physics is inactive at scale.
TEMP, PRCP, DH, RHMEAN	Daymet v4 (1 km gridded)	ORNL DAAC single-pixel API; open (CC0-equivalent for U.S. government / DAAC terms)	TEMP = mean of tmax/tmin; DH = incident solar radiation (srad, a physical radiative index replacing the original ambiguous DH, on a different scale); RHMEAN derived from vapour pressure via the Tetens saturation relation.
WDSP	gridMET	Climatology Lab / Northwest Knowledge Network THREDDS NetCDF-Subset Service; open	daily mean wind speed at the station coordinate.

Forty stream gages across 17 U.S. states (≥ 55 % water-temperature coverage over 2006–2020) span free-flowing and dam-regulated rivers; water-temperature coverage ranges 0.56–1.00 (median ≈ 0.85). The two assembled panels (`data_usgs/panel_usgs.parquet`, six variables; `data_usgs/panel_usgs_wind.parquet`, seven variables) and station metadata (`data_usgs/stations_meta.csv`, with USGS site numbers and coordinates) are included so the experiments reproduce without re-downloading; the raw per-site downloads can be regenerated by re-running the acquisition script. All three sources are public domain or open-licensed.

Code availability and reproducibility

All code, the fixed train/validation/calibration/blind-test boundaries, unit tests (leakage, splits, metrics, con-formal), and one-command reproduction (`scripts/run_all.sh`) are in this repository. Per-day predictions, model weights and logs are regenerable by the staged scripts (`scripts/01–13`). An adversarial internal review (multiple independent expert lenses) checked every headline claim against the result ta-

bles; all negative results and ablations are retained, and a finding-by-finding disposition is provided in outputs/reports/review_response.md.

References

Formatted following the Journal of Hydrology house style (Elsevier numeric-author-year). Cross-reference labels in the manuscript text use the F-numbers retained for traceability.

Stream temperature: reviews, baselines and learned models

Corona, C. R., Hogue, T. S., 2025. Machine learning in stream and river water temperature modeling: a review and metrics for evaluation. *Hydrology and Earth System Sciences* 29, 2521–2549. <https://doi.org/10.5194/hess-29-2521-2025>. [F1]

Feigl, M., Lebedzinski, K., Herrnegger, M., Schulz, K., 2021. Machine-learning methods for stream water temperature prediction. *Hydrology and Earth System Sciences* 25, 2951–2977. <https://doi.org/10.5194/hess-25-2951-2021>. [F2]

Zwart, J. A., Sadler, J. M., Appling, A. P., Read, J. S., Oliver, S. K., Diaz-Gonzalez, M., Bertassello, L. E., Jia, X., 2023. Evaluating deep learning architecture and data assimilation for improving water temperature forecasts at unmonitored locations. *Frontiers in Water* 5, 1184992. <https://doi.org/10.3389/frwa.2023.1184992>. [F3]

Jia, X., Zwart, J., Sadler, J., Appling, A., Oliver, S., Markstrom, S., Willard, J., Xu, S., Steinbach, M., Read, J., Kumar, V., 2021. Physics-guided recurrent graph model for predicting flow and temperature in river networks, in: *Proceedings of the 2021 SIAM International Conference on Data Mining (SDM)*. SIAM, pp. 612–620. <https://doi.org/10.1137/1.9781611976700.69> (preprint: arXiv:2009.12575). [F4]

Luo, S., Yu, R., Chen, S., Jia, X., Kumar, V., 2025. Geo-aware models for stream temperature prediction across different spatial regions and scales, in: *Proceedings of the 33rd ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*. ACM. <https://doi.org/10.1145/3748636.3762716> (preprint: arXiv:2510.09500). [F5]

Rahmani, F., Lawson, K., Ouyang, W., Appling, A., Oliver, S., Shen, C., 2021. Exploring the exceptional performance of a deep learning stream temperature model and the value of streamflow data. *Environmental Research Letters* 16, 024025. <https://doi.org/10.1088/1748-9326/abd501>. [F7]

Rahmani, F., Shen, C., Oliver, S., Lawson, K., Appling, A., 2021. Deep learning approaches for improving prediction of daily stream temperature in data-scarce, unmonitored, and dammed basins. *Hydrological Processes* 35, e14400. <https://doi.org/10.1002/hyp.14400>. [F8]

Rahmani, F., Appling, A., Feng, D., Lawson, K., Shen, C., 2023. Identifying structural priors in a hybrid differentiable model for stream water temperature modeling. *Water Resources Research* 59, e2023WR034420. <https://doi.org/10.1029/2023WR034420>. [F9]

Sadler, J. M., Appling, A. P., Read, J. S., Oliver, S. K., Jia, X., Zwart, J. A., Kumar, V., 2022. Multi-task deep learning of daily streamflow and water temperature. *Water Resources Research* 58, e2021WR030138. <https://doi.org/10.1029/2021WR030138>. [F10]

Topp, S. N., Appling, A. P., Diaz-Gonzalez, M. A., Heinrich, L. R., Zwart, J. A., Read, J. S., Oliver, S. K., 2023. Stream temperature predictions for river basin management in the Pacific Northwest and Mid-Atlantic regions using machine learning. *Water Resources Research* 59, e2022WR033880. <https://doi.org/10.1029/2022WR033880>. [F11]

Toffolon, M., Piccolroaz, S., 2015. A hybrid model for river water temperature as a function of air temperature and discharge. *Environmental Research Letters* 10, 114011. <https://doi.org/10.1088/1748-9326/10/11/114011>. [F12]

Piccolroaz, S., Calamita, E., Majone, B., Gallice, A., Siviglia, A., Toffolon, M., 2016. Prediction of river water temperature: a comparison between a new family of hybrid models and statistical approaches. *Hydrological Processes* 30, 3901–3917. <https://doi.org/10.1002/hyp.10913>.

Mohseni, O., Stefan, H. G., Erickson, T. R., 1998. A nonlinear regression model for weekly stream temperatures. *Water Resources Research* 34, 2685–2692. <https://doi.org/10.1029/98WR01877>.

Caissie, D., 2006. The thermal regime of rivers: a review. *Freshwater Biology* 51, 1389–1406. <https://doi.org/10.1111/j.1365-2427.2006.01597.x>.

Methods — machine learning, calibration, statistics

Romano, Y., Patterson, E., Candès, E. J., 2019. Conformalized quantile regression, in: *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*, pp. 3543–3553.

Vovk, V., Gammerman, A., Shafer, G., 2005. *Algorithmic Learning in a Random World*. Springer, New York.

Martins, A. F. T., Astudillo, R. F., 2016. From softmax to sparsemax: a sparse model of attention and multi-label classification, in: *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, pp. 1614–1623.

Bai, S., Kolter, J. Z., Koltun, V., 2018. An empirical evaluation of generic convolutional and recurrent networks for sequence modeling. *arXiv:1803.01271*.

Shazeer, N., Mirhoseini, A., Maziarz, K., Davis, A., Le, Q., Hinton, G., Dean, J., 2017. Outrageously large neural networks: the sparsely-gated mixture-of-experts layer, in: *International Conference on Learning Representations (ICLR)*.

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., Liu, T.-Y., 2017. LightGBM: a highly efficient gradient boosting decision tree, in: *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, pp. 3146–3154.

Diebold, F. X., Mariano, R. S., 1995. Comparing predictive accuracy. *Journal of Business & Economic Statistics* 13, 253–263. <https://doi.org/10.1080/07350015.1995.10524599>.

Harvey, D., Leybourne, S., Newbold, P., 1997. Testing the equality of prediction mean squared errors. *International Journal of Forecasting* 13, 281–291. [https://doi.org/10.1016/S0169-2070\(96\)00719-4](https://doi.org/10.1016/S0169-2070(96)00719-4).

Künsch, H. R., 1989. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics* 17, 1217–1241. <https://doi.org/10.1214/aos/1176347265>.

Wilks, D. S., 2011. *Statistical Methods in the Atmospheric Sciences*, 3rd ed. Academic Press, Oxford. (Chapters on bootstrap CIs, Brier/Brier-skill, AUPRC, and the Relative Economic Value of probabilistic forecasts.)

Richardson, D. S., 2000. Skill and relative economic value of the ECMWF ensemble prediction system. *Quarterly Journal of the Royal Meteorological Society* 126, 649–667. <https://doi.org/10.1002/qj.49712656313>.

Gneiting, T., Raftery, A. E., 2007. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association* 102, 359–378. <https://doi.org/10.1198/016214506000001437>.

Kingma, D. P., Ba, J., 2015. Adam: a method for stochastic optimization, in: *International Conference on Learning Representations (ICLR)*.

Loshchilov, I., Hutter, F., 2019. Decoupled weight decay regularization, in: *International Conference on Learning Representations (ICLR)*.

Nash, J. E., Sutcliffe, J. V., 1970. River flow forecasting through conceptual models part I — a discussion of principles. *Journal of Hydrology* 10, 282–290. [https://doi.org/10.1016/0022-1694\(70\)90255-6](https://doi.org/10.1016/0022-1694(70)90255-6).

Gupta, H. V., Kling, H., Yilmaz, K. K., Martinez, G. F., 2009. Decomposition of the mean squared error and NSE performance criteria: implications for improving hydrological modelling. *Journal of Hydrology* 377, 80–91. <https://doi.org/10.1016/j.jhydro.2009.08.003>.

Software and data sources

De Cicco, L. A., Lorenz, D., Hirsch, R. M., Watkins, W., Johnson, M., 2024. dataRetrieval: R packages for discovering and retrieving water data available from federal hydrologic web services. U.S. Geological Survey. <https://doi.org/10.5066/P9X4L3GE>. [D1]

Hodson, T. O., Decker, J. K., 2023. dataretrieval (Python): a Python package for discovering and retrieving water data available from U.S. federal hydrologic web services. U.S. Geological Survey. <https://doi.org/10.5066/P94I5TX3>.

U.S. Geological Survey, 2024. National Water Information System (NWIS). Web services accessed via dataRetrieval; parameter codes 00010 (water temperature), 00060 (discharge), 00065 (gage height). <https://waterdata.usgs.gov/nwis>. [public domain]

Thornton, P. E., Shrestha, R., Thornton, M., Kao, S.-C., Wei, Y., Wilson, B. E., 2022. Daymet: daily surface weather data on a 1-km grid for North America, version 4 R1. ORNL DAAC, Oak Ridge, Tennessee, USA. <https://doi.org/10.3334/ORNLDAAAC/2129>. [D2]

Abatzoglou, J. T., 2013. Development of gridded surface meteorological data for ecological applications and modelling. *International Journal of Climatology* 33, 121–131. <https://doi.org/10.1002/joc.3413>. (gridMET / METDATA.) [D3]

Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., Bai, J., Chintala, S., 2019. PyTorch: an imperative style, high-performance deep learning library, in: *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*, pp. 8026–8037.

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, É., 2011. Scikit-learn: machine learning in Python. *Journal of Machine Learning Research* 12, 2825–2830.

Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., Burovski, E., Peterson, P., Weckesser, W., Bright, J., van der Walt, S. J., Brett, M., Wilson, J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson, E., Carey, C. J., Polat, I., Feng, Y., Moore, E. W., VanderPlas, J., Laxalde, D., Perktold, J., Cimrman, R., Henriksen, I., Quintero, E. A., Harris, C. R., Archibald, A. M., Ribeiro, A. H., Pedregosa, F., van Mulbregt, P., SciPy 1.0 Contributors, 2020. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods* 17, 261–272. <https://doi.org/10.1038/s41592-019-0686-5>.